

Machine Learning Engineer Nanodegree

Capstone Project

Tom Martin 30th January 2017

I. Definition

Project Overview

My project will examine the MNIST database of handwritten digits. The project aims are to deduce an accurate SVM classifier for this dataset and analyse its performance against a benchmark study, noted below.

This is a very well known dataset having attracted a great deal of academic attention since its inception. More broadly, the analysis of automated handwriting recognition has applications for fields where it is important to quickly and securely process handwritten documents at scale. For instance, this can be useful in processing historical documents, input to handheld devices via a stylus or pen, or determining authorship of incriminating documents.

Over the years, it has proved fruitful territory for examining a range of machine learning classifiers, such as linear classifiers[1], svm[2], k-nearest neighbours[3], and a range of neural network implementations[4]. There have therefore been a number of different approaches shown to be suitable to classify the dataset correctly. A paper by Hartwick[5] is particularly relevant for this project as he has provided a clear analysis of the dataset without any further preprocessing with an SVM classifier. For these reasons, this paper will form a benchmark for the following analysis.

Problem Statement

The capstone will attempt to train and tune a SVM classifier that is able to correctly determine the number intended from the supplied image of a handwritten sample. The model produced will be trained, validated and tested against the supplied dataset. The success of the classifier will be measured using the SciKit-Learn metric's module `metrics.accuracy_score`, `metrics.confusion_matrix`. From these metrics we can derive both the error rate and per digit error rate in order to enable a direct comparison with the benchmark model below. It should be mentioned that some of these metrics are typically used in problems of binary classification, but can be generalised for an arbitrary number of classes[6]. This is covered in more detail in the "Metrics" section.

I propose using a SVM classifier to train a solution that, with a reasonable level of accuracy, correctly maps a handwritten sample to the correct digit. A

supervised classifier should be an appropriate solution to the problem as we have training data. There are also a number of academic studies that have had success with SVM classifiers[2]. Before building the model, I will use principal component analysis (PCA) with dimension reduction. PCA is used to reduce the feature space of the training data to reduce the overall training and testing time. This is especially useful for a dataset both as large and feature-rich as the MNIST dataset. It has been shown that in general PCA does not negatively impact the accuracy of a classifier, and has even been shown to boost the accuracy of the SVM classifier[7]. I will evaluate the classifier using two different kernels: Gaussian, or RBF kernel, and polynomial. The choice is for two reasons: Firstly, Both these kernels are useful in cases such as this where the data set is not linearly separable. Secondly, I want to make a direct comparison with the benchmark study, which used the Gaussian kernel.

The trained classifier can be evaluated using a confusion matrix, and derived metrics to determine its degree of success. Furthermore, we can consider a number of previous models[8] of the dataset using a SVM classifier, which have accuracies around 99%.

Metrics

The evaluation metrics for this model will be the confusion matrix, accuracy, error rater, and per digit error rate. The latter two are to enable a direct comparison with the benchmark, which calculates these values in the paper.

I will be focussing on the confusion matrix and accuracy to evaluate the classifier for the two main reasons. Firstly, the labels in the dataset are fairly uniformly distributed, as shown in figure 2 below. This means we can take the simpler option of just using accuracy as we do not need to consider imbalances between classes - accuracy will strongly correspond to other measures such as precision in this case. Secondly, this is a multi-class classification problem where we are more interested in correct classifications than misclassifications, so accuracy is sufficient to meaningfully evaluate the classifier on its own. The confusion matrix will still contain useful information especially if there is strong presence of misclassifications on a per digit bases, which would not be as clear from accuracy alone.

A supervised classifier such as SVM has known labelled data, so we can determine the number of true positives, true negatives, false positives, and false negatives these are ultimately derived from the confusion matrix. To be clear, these terms are defined as follows:

- True positives: Entries that are correctly labelled
- True negatives: Entries that are correctly rejected
- False positives: Entries that a wrongly identified with a given label
- False negatives: Entries for a given label that are wrongly identified with other labels

In the general case, a confusion matrix is simply a matrix illustrating the mapping from the true labels to the predicted labels. Elements along the diagonal represent a correct classification, whereas the off-diagonal represent a misclassification. A confusion matrix can be a useful check to see what digits in particular are most likely confused for one another. From here we can derive the accuracy.

Accuracy is given by the total number of correct classifications, both true positives and true negatives divided by the total dataset population. This can be given by the following equation,

$$\text{accuracy} = (\text{tp} + \text{tn}) / (\text{tp} + \text{tn} + \text{fp} + \text{fn})$$

where tp, tn, fp, and fn stand for true positive, true negative, false positive, false negative respectively.

From this, the error rate or rate at which the classifier misclassifies can be derived,

$$\begin{aligned} \text{error rate} &= 1 - \text{accuracy} \\ &= 1 - (\text{tp} + \text{tn}) / (\text{tp} + \text{tn} + \text{fp} + \text{fn}) \\ &= (\text{fp} + \text{fn}) / (\text{tp} + \text{tn} + \text{fp} + \text{fn}) \end{aligned}$$

To determine the error rate for a given digit I refer to the benchmark paper[5], which defines the error rate as the ratio of false negatives to the sum of true positives and false negatives.

$$\text{error rate (for given digit)} = \text{fn} / \text{tp} + \text{fn}$$

This can be derived from the confusion matrix: For a given row, divide the off diagonal entries by the sum of all entries for that row.

These metrics altogether will give us a means to determine how well the classifier correctly labels the digits as well the error rate per digit. The error rate as well as all the other metrics discussed in this section will be calculated using the SciKit-Learn `metrics.accuracy_score`, `metrics.confusion_matrix` methods.

II. Analysis

Data Exploration

The MNIST dataset contains 70000 samples of handwritten digits, labelled from 0 to 9. These are split into subsamples of 60000 and 10000 for training and testing respectively. The samples themselves have been centred and normalised to a grid size of 28-by-28 pixels, with each training entry composed of 784 features, corresponding to the greyscale level for each pixel. The MNIST dataset in this case will be the MNIST original[9] dataset obtained via the mldata repository using SciKit-Learn's `datasets.fetch_mldata` method. A sample of the dataset is given below.



Figure 1: Sample of MNIST Dataset

The class labels in the testing set are roughly uniformly distributed, with the number of occurrences of each label ranging from around 6300 and 7900. The distribution is shown graphically in figure 2 below. This distribution of labels means that no special sampling needs to take place to train and test correctly, and as discussed above, means we can opt for the simpler choice of deriving only the accuracy as the evaluation metric.

On a historical note, this dataset is the result of subsampling the original NIST dataset so that it was overall more consistent, and more suitable for machine learning: mixing together the original training and testing sets. The samples were collected from a combination of American Census Bureau employees and American high school students. This dataset contains both training and testing samples, so no further data is needed to evaluate the classifier.

Exploratory Visualization

The plot below illustrates the distribution of labels in the target data, showing them to be fairly uniformly distributed. This is useful in relation to the choice of evaluation metric, as noted above. Due to this uniformity, we can opt for the simpler choice of using just the accuracy as we do not need to consider imbalances and skew.




Figure 2: Frequency of Occurrence of Class Labels

Algorithms and Techniques

This report will train a scalable vector machine (SVM) classifier to correctly label the samples. SVM is a way of classifying data sets with a hyperplane or boundary between data point clusters. This is achieved by algorithmically finding the hyperplane with maximum margin, or in other words finding the line that maximises the distance between itself and the points of the data set. SVM is particularly powerful when dealing with nonlinear data by employing the so-called kernel trick: mapping the input data points to a higher dimensional feature space.

The SVM provided by SciKit-Learn's `svm.SVC` module supports the following parameters relevant for this investigation,

- C - penalty parameter, how closely hyperplane follows the datapoints
- kernel - the kernel choice, either RBF or polynomial
- degree - the degree of the polynomial kernel
- gamma - kernel coefficient, defines radius of influence for a given datapoint

The specific parameter values will be decided by cycling through a range of values starting from a set of values following the relevant literature. This technique is a simple and effective means of optimising the classifier, as a way of justifying our initial assumptions. Other parameters not specified here will be used at their default values.

Due to the size of the dataset used in this project it is necessary to use a hybrid approach[10]. This means initially running the dataset through a k nearest neighbours classifier (KNN) and only run digits through the SVM classifier where the k neighbours do not have the same label. Together this known as KNN-SVM, and has the dual advantage of achieving a high level of accuracy on a feature rich dataset using SVM, whilst offloading a large proportion of the classification problem to a much cheaper KNN classifier. In addition to the SVM parameters I will vary the following,

- PCA: `n_components`
- KNN: `n_neighbors`

to determine their effect on the classifier accuracy. The classifier will be trained using a validation set by performing a test/train split. Each trained classifier will be used to evaluate the test set. The performance metrics to be used are discussed above.

Preprocessing will be done using PCA with dimension reduction. This is discussed at greater length below.

Benchmark

This dataset is very well studied and as such, there are many comparable studies to check against. For the project, I will make direct comparison to the paper referenced above by Hartwick[5]. This paper produces results for a SVM classifier with Guassian kernel, with parameters,

```
C = 10^6  
gamma = (1/len(features)) * 10^-3.5 (approx. 4 * 10^-7)
```

The model in this paper achieves a error of around 1.4% against the MNIST testing set, the paper also gives the per digit error rate. Given all these results and the availability of the identical testing set, a direct comparison with this paper's results is possible.

However, unlike the benchmark study, I will perform preprocessing on the dataset using PCA. Following the study by Lei and Govindaraju[7] I will choose to model with several different numbers of principal components between 25 and 100, as this is where they found a boosted classifier performance. I will also use a hybrid approach, using a combination of KNN and SVM classifiers, due to the size of the dataset as discussed above.

III. Methodology

Data Preprocessing

The data was fetched using SciKit-Learn's `datasets.fetch_mldata`, which creates a local cache for subsequent reads.

Data preprocessing was achieved using PCA with dimensional reduction. This was to enable a speed up in the training and execution time of the classifier, which has been shown to not degrade performance in general - and has even let to some improvement. This is also a very necessary step when processing such a large and feature reach dataset such as MNIST especially when training on general purpose hardware. The number of components to use in this step comes follows from two considerations. Firstly, a previous study using PCA on the MNIST dataset showed boosted SVM performance for a number of principal components below 100[7]. Secondly, by looking at the explained variance ratio, we see that over 91% of the variance is accounted for by the top 100 components.

```
pca = PCA().fit(X) # where X is the imported training set  
explained_variance = pca.explained_variance_ratio_  
sum(explained_variance[101]) # => 91.6%
```

For my implementation, I used the Python `pickle` module to serialise the target and and preprocessed data samples to easily share across files.

PCA was performed before the test train split to ensure consistent analysis, the dimensionality of the training data should be the same as the test data.

See `Downloading PCA and Caching.ipynb` in the `code_samples` directory for implementation details.

Implementation

The full implementation for this project is illustrated in the figure below. As a whole, the project follows the approach given by Zhang et al.[10] known as SVM-KNN. I was not aware of the need to use a hybrid approach when I initially planned this project but this SVM Implementation does not scale well for larger datasets, with a reported order of growth between the square of the number of samples and the cube the number of samples[11].

Going from top to bottom through the above figure. Given the initial dataset, I performed PCA with dimensional reduction for three different numbers of principal components and serialised the result with the `pickle` module. A test/train split was then performed before moving on to training, in order to have a validation set to train against. For a given preprocessed dataset, I ran this though the KNN classifier to perform an initial classification. A subsequent call to the SVM classifier was only made if the given k neighbours do not have the same label i.e. a further classification is required. It is at this point that the distance matrix is converted to a kernel matrix for SVM classifier.

All the algorithms and modules used in the project implementation come from SciKit-Learn 0.18.1 library. Those being,

for classification,

- `decomposition.PCA`
- `neighbors.KNeighborsClassifier`
- `svm.SVC`

for test/train split

- `model_selection.train_test_split`

and to derive evaluation metrics

- `metrics.accuracy_score`
- `metrics.confusion_matrix`

The training and evaluation step was done over a set of variables for both the KNN and SVM classifiers to perform a basic hyperparameter optimisation.

See the notebooks in the `code_samples` directory for implementation details.




Figure 3: Project Implementation

Refinement

The classifier refinement step is bundled as part of the classifier training. I used a test/train split to produce a validation set to use as part of the training step, then run the trained classifier on the test set. I used a range of parameters following from the my reading of the literature, this meant varying the number of principal components, the number of nearest neighbours, and the SVM parameters C , γ , kernel and degree - for the polynomial kernel only.

A subset of results for the best performing parameters are shown below.

Varying nearest neighbours, k , for constant number of principal components

k	PC	Parameter Set	Test Set Accuracy
50	100	1	0.9843
100	100	1	0.9858
200	100	1	0.9858

Varying number of principal components for constant nearest neighbours k

k	PC	Parameter Set	Test Set Accuracy
100	25	2	0.9822
100	50	1	0.9849
100	100	1	0.9858

Parameter sets

Parameter set 1, all eight combinations of the following,

- $C = \{ 1, 10, 100, 1000 \}$
- $\text{degree} = 2$
- $\gamma = (10^{-3}, 10^{-4})$
- $\text{kernel} = \text{polynomial}$

Parameter set 2, either the two combinations of,

- $C = 1$
- $\text{degree} = 2$
- $\gamma = (10^{-3}, 10^{-4})$
- $\text{kernel} = \text{polynomial}$

or the eight combinations of,

- $C = \{ 1, 10, 100, 1000 \}$
- $\text{degree} = 3$
- $\gamma = (10^{-3}, 10^{-4})$

- `kernel = polynomial`

Although varying the number of parameters did not appreciably increase the testing accuracy, we can spot a few trends,

- Increasing `k` improved the classifier accuracy, at least up to a point - see rows where `k` equals 100 and 200 in the first table. However, this comes at a cost of longer processing times
- Increasing `PC` improved classifier accuracy, more significantly than `k`. This is perhaps to be expected as the resulting dataset has more features to use. However, again this comes at a cost of longer processing times
- The parameter selection for SVM was not very significant for the final classifier - the best performing classifiers used eight parameter configurations. This is probably a consequence of using a hybrid approach: The input space to the SVM is fairly small
- A low degree polynomial kernel outperformed all other kernel choices - this is common in the NLP space, as these least overfit the data

See `Train and Optimise Classifier.ipynb` in the `code_samples` directory for implementation details.

IV. Results

Model Evaluation and Validation

The best performing classifier had the following parameters,

- `k = 100`
- `PC = 100`
- SVM parameters are any from parameter set 1, above

These parameters were selected by comparing each trained classifier against the test set. These parameters make sense in light of theoretical considerations,

- The number of nearest neighbours, `k`, should be approximately $\sqrt{N}/2$
- The number of principal components account for 91% of the variance, see the “Data Preprocessing” section above

See `Train and Optimise Classifier.ipynb` and `Evaluate Classifier.ipynb` in the `code_samples` directory for implementation details.

Justification

The final classifier used an arbitrary selection of parameters from parameter set one.

The benchmark reported above produced a SVM classifier with an error of 1.4%. This compares favourably with the classifier trained in this project with has an

error of 1.42% (corresponding to an accuracy of 98.58%). Discrepancies between the benchmark results may be due to,

- Preprocessing with PCA on the initial dataset, the benchmark study did not use any preprocessing
- Combination of KNN and SVM classifiers, the benchmark study trained the SVM directly on the MNIST set
- Different parameters used in final classifier, potentially related to previous two points

More detailed comparison with the benchmark follows by comparing the error per digit. This gives some indication of how the two classifiers differ and highlight the effect of preprocessing. Accounting for the overall larger error for the classifier trained in this project in relative terms, the error per digit was fairly similar despite the preprocessing. In both cases the digits 1, 0, 6 has the lowest error rates, and digit 9 had significantly greater error rates. Interesting differences between the classifiers are highlighted by looking at the error rate for digit 4, which had the 4th lowest error rate in the benchmark, but had the 3rd worst error rate in my trained classifier. It is not clear how this relates to PCA and KNN.

Digit	Benchmark Error	Error per Digit
0	0.0061	0.0022
1	0.0052	0.0053
2	0.0164	0.0104
3	0.0148	0.0176
4	0.0132	0.0214
5	0.0168	0.0232
6	0.0104	0.0101
7	0.0184	0.0139
8	0.0184	0.0136
9	0.0277	0.0257

V. Conclusion

Free-Form Visualization

The results for this project can be powerfully illustrated using a confusion matrix shown below. This is from the results of the optimised classifier detailed in the section “Justification” above.

This demonstrates a very powerful performance overall from the classifier, as well as highlighting the digits that proved more difficult relative to others. This is discussed in more detail above.




Figure 4: Predicted vs True label

Reflection

The workflow for this project was as follows,

1. Research of problem area and relevant benchmark in the literature
2. Finding and downloading relevant dataset
3. Preprocessed dataset using PCA
4. Initial classification of dataset using KNN
5. Trained SVM using range of parameters - only used samples not correctly labelled by KNN
6. Produced more results for most successful KNN-SVM classifier

The most difficult stages in the project were also happily the most interesting, those being steps 4 and 5. I was not clear in my previous research and project proposal how to efficiently run the SVM classifier against such a large dataset. This proved to be a real bottleneck for the project and required further investigation to proceed. However the hybrid approach used was both very effective and also a first for me.

The final model is robust, effective and efficient. In this project I have been able to derive a classifier with over 98% accuracy, highly competitive with the benchmark study, in a total running time of under 160s (on my laptop). This is extremely promising moving into a more general setting.

Improvement

The major bottleneck for this project was the implementation of the SVM classifier used by SciKit-Learn. This SVM does not scale well for larger datasets as noted above in the “Implementation” section. A number of solutions to explore include trying to better utilise caching of calculated distances in memory, trying alternatives to SciKit-Learn like Vowpal Wabbit, or make better use of dedicated hardware, such as running the SVM on the GPU rather than a general purpose CPU. If a significant performance improvement can be made for the SVM classifier then I believe a future classifier could outperform the model created in this project. This is because a more optimised classifier would be less reliant on preprocessing or hybrid approaches.

References

- [1] <http://yann.lecun.com/exdb/publis/pdf/lecun-98.pdf> “Gradient-Based Learning Applied to Document Recognition”
- [2] <https://people.eecs.berkeley.edu/~malik/cs294/decoste-scholkopf.pdf> “Training Invariant Support Vector Machines”

- [3] <http://citeseerx.ist.psu.edu/viewdoc/download;jsessionid=E0F3BDC7642FBA1D8E2811526BD0E596?doi=10.1.1.106.3963&rep=rep1&type=pdf> “Deformation Models for Image Recognition”
- [4] <https://www.microsoft.com/en-us/research/publication/best-practices-for-convolutional-neural-networks-applied-to-visual-document-analysis/> “Best Practices for Convolutional Neural Networks Applied to Visual Document Analysis”
- [5] <https://cseweb.ucsd.edu/~jmcauley/cse190/reports/fa15/025.pdf> “Reproducing Results of Guassian Kernel SVM classifiers on the MNIST Dataset”
- [6] <http://softclassval.r-forge.r-project.org/2013/2013-01-03-ChemomIntellLabSystemTheorypaper.html> “Validation of Soft Classification Models using Partial Class Memberships: An Extended Concept of Sensitivity & Co. applied to Grading of Astrocytoma Tissues”
- [7] https://www.researchgate.net/profile/Giovanni_Felici/publication/226795010_Feature_Selection_for_Data_Mining/links/53e413d70cf25d674e94b475.pdf#page=78 “Speeding Up Multi-class SVM Evaluation by PCA and Feature Selection”
- [8] <http://yann.lecun.com/exdb/mnist/> “The MNIST Database of Handwritten Digits”
- [9] <http://mldata.org/repository/data/viewslug/mnist-original/> “MNIST (original)”
- [10] https://www.vision.caltech.edu/Image_Datasets/Caltech101/nhz_cvpr06.pdf SVM-KNN: Discriminative Nearest Neighbor Classification for Visual Category Recognition
- [11] <http://scikit-learn.org/stable/modules/svm.html#complexity> “Support Vector Machines - Complexity”